

COSC3500: High Performance Computing

Semester 2, 2021

n-body Gravity Simulation

Project Report - Milestone 2

by Tom Stephen
45811449

Contents

1	Serial	2
1.1	Introduction	2
1.2	Background	2
1.3	Implementation	3
1.3.1	Algorithm	3
1.3.2	Program Design	4
1.3.3	Profiling and Optimisation	4
1.4	Verification	6
1.5	Performance Results	8
2	Parallel	10
2.1	Parallelisation Strategies	10
2.2	CUDA Implementation	10
2.2.1	Profiling	12
2.3	Verification	12
2.4	Experimentation	14
2.5	Conclusion	17
3	Appendix	19
3.1	Serial Bash Scripts	19
3.1.1	Serial Validation	19
3.1.2	Parallel Validations	20
3.2	Performance Analysis	22
3.2.1	Serial Flags	22
3.2.2	Parallel Flags	22
3.2.3	Misc	23
3.3	Main Source Code	24
3.3.1	Serial	24
3.3.2	Parallel	27
3.3.3	Helper functions	33

Chapter 1

Serial

1.1 Introduction

This project attempts to implement and later optimise an interacting particle system - specifically, gravitational interactions between massive bodies in space - using high performance computing methods. This chapter introduces the theory required, motivates this task, then demonstrates and benchmarks a serial implementation. Some basic analysis is conducted to verify that the implementation is performing correctly, and the time dependance on input parameters is presented. The next chapter is focused on how this implementation may be optimised using parallel computing techniques.

1.2 Background

Increasingly, the space surrounding Earth has become littered by satellites and “space junk” that any new space endeavor must avoid. Should this be an isolated system, where each satellite/space junk/object (henceforth referred to as a “body”) only experiences an attractive force with the Earth, it would be relatively trivial to map the positions over time for all bodies analytically. However, gravitational interactions with the sun, moon, asteroids and other pieces of debris have a non-trivial effect on each other, and thus must be taken into account for a truly accurate predication of future positions. In any useful model, we would wish to track as many bodies as possible. This leads to performance difficulties as the force each body imparts on every other body must be calculated.

The gravitational attractive force of object 1 acting on object 2 is given by

$$F = G \frac{m_1 m_2}{r^2} \quad (1.1)$$

where m_1, m_2 are the masses of the interacting objects, and r is the distance between them. G is the gravitational constant, and is $G = 6.67 \times 10^{-11} \text{Nm}^2/\text{kg}^2$ [1]. Conveniently, the force on object 1 from object 2 is equal to the force of object 2 from object 1, so we may cut down our computations. Once we know the total force acting on an object, F_{total} , we know the acceleration given by newtons second law of motion,

$$a = \frac{F_{\text{total}}}{m} \quad (1.2)$$

Eventually, we will wish to initialise an example solar system, comprising of a supermassive body (i.e, a star), orbited by some number of much smaller bodies (i.e., planets). The velocity of the orbiting bodies must be perpendicular to the main acceleration, i.e., towards the supermassive body, and the velocity must be such that the gravitational acceleration is equal to the centripetal acceleration. With supermassive body having mass M , orbiting body with mass m , distance r

and velocity v , we have the gravitational acceleration as

$$a = \frac{GM}{r^2} \quad (1.3)$$

[1] centripetal acceleration,

$$a = \frac{v^2}{r} \quad (1.4)$$

[1] Equating and solving for v ,

$$\frac{v^2}{r} = \frac{GM}{r^2} \quad (1.5)$$

$$v_{\text{orbit}} = \sqrt{\frac{GM}{r}} \quad (1.6)$$

So, given some arbitrary r , we may calculate the orbital velocity.

1.3 Implementation

1.3.1 Algorithm

All tests in this report use a time range of 100s, with a step size of $h = 0.1$ so consistency. The number of bodies, n , is varied for each test.

Algorithm 1 Main Loop Algorithm

```
create  $n$ -bodyarrays
initialise  $n$ -bodies
simulate system
delete  $n$ -bodyarrays
```

A separate array is used for x position, y position, x, y velocities, x, y accelerations and body mass, where each index represents a unique body. It is tempting to structure the project using a body class or struct (which, in C++, is essentially a class), and then store the bodies in a dynamic vector. However, this adds a few levels of overhead - having to access classes is significantly slower than simply indexing an array. All calculations are done in-place, to reduce memory usage. Similarly, if selected, the current state is output after each time step, so that the arrays need only contain the current state, and not be a matrix including all previous states.

The function to initialise the bodies is not particularly complex. It simply sets the first body to be the “star”, then loops through each of the remaining bodies, generating a random distance and angle from the center. The position is then calculated from this, and the orbital speed is calculated using equation 1.6.

The simulation function is also quite simple. At each time step, it calls the `time_step` function, and then outputs some data. The `time_step` follows the following algorithm,

Algorithm 2 Time Step Function

```
call update_acceleration
for each body  $\in$  all bodies do
     $v_x \leftarrow v_x + h \cdot a_x$ 
     $v_y \leftarrow v_y + h \cdot a_y$ 
     $x \leftarrow x + h \cdot v_x$ 
     $y \leftarrow y + h \cdot v_y$ 
end for
```

and so, this is a simply implementation of Euler's method. Euler's method is arguably the most simple iterative method, whereby one updates the value of some function, say $f(x_n)$ to $f(x_{n+1})$ by adding the product of the derivative of the function, $f'(x_n)$ and the difference in the points, $x_{n+1} - x_n$. That is,

$$f(x + h) = f(x) + h * f'(x)$$

. As h becomes small, the updates each time are smaller and in turn more accurate - they approximate less of the acceleration and velocity changes.

Finally, the `update_acceleration` function. This is the most computationally intensive function by far. It looks at every pair of body to calculate the gravitational force that each exerts on one another, and then calculates accelerations experienced by each body, and appends that to the acceleration vectors. It's algorithm is as follows,

Algorithm 3 Update Acceleration Function

```

for  $i \in 1 : n$  do                                     ▷ Set all accelerations to 0
     $a_x[i] \leftarrow 0$ 
     $a_y[i] \leftarrow 0$ 
end for
for  $i \in 1 : n$  do
    for  $j \in (i + 1) : n$  do
         $\text{dist} = \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}$            ▷ Calculate distance between bodies
         $F_x = -G \cdot m_i \cdot m_j \cdot (x_i - x_j) \cdot \text{dist}^{-3}$ 
         $F_y = -G \cdot m_i \cdot m_j \cdot (y_i - y_j) \cdot \text{dist}^{-3}$ 
         $a_x[i] \leftarrow a_x[i] + F_x/m_i, \quad a_x[j] \leftarrow a_x[j] - F_x/m_j$ 
         $a_y[i] \leftarrow a_y[i] + F_y/m_i, \quad a_y[j] \leftarrow a_y[j] - F_y/m_j$ 
    end for
end for

```

The first for loop only takes $\mathcal{O}(n)$ time, so does not contribute so much to the run time. The second for loop, however, runs $\mathcal{O}(n!)$ times! The operation inside the for loop are also relatively slow - particularly the square root and repeated floating point division.

1.3.2 Program Design

Most of the design choices are discussed above as they have been made for optimisation purposes. The project has been created in C++11, making use of a few built in libraries. C++ itself is used over C mostly for being able to handle command line arguments - there is no particular reason it couldn't be ported to C. Plotting is done in gnuplot, and animation is done with a Python script. Details for these are below.

1.3.3 Profiling and Optimisation

To determine which functions where the slowest in runtime, the C++ build command was run with the `-pg` flag, and without any optimisation flags. `gprof` was then used to produce a profile of the run time. The first few lines of the analysis are attached below, with simulation set up of 20 bodies, timestep 1s, final time 100,000s:

Each sample counts as 0.01 seconds.

%	cumulative	self	self	self	total	name
time	seconds	seconds	calls	ms/call	ms/call	
93.42	0.70	0.70	100001	0.01	0.01	update_acceleration(double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*)
2.67	0.72	0.02	400004	0.00	0.00	euler(double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*)
2.67	0.74	0.02	100001	0.00	0.01	time_step(double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*)
1.33	0.75	0.01	100001	0.00	0.00	kinetic_energy(double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*)
0.00	0.75	0.00	100001	0.00	0.00	write_state(double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*, double*)

```

0.00      0.75      0.00      2496      0.00      0.00  std::__detail::_Mod<unsigned long, 429496
...

```

Clearly, a majority of the run time is spent in the `update_acceleration` function, even though it has fewer calls than, for example, the `euler` function. Clearly, the bulk of the optimisation should be spent here. Since the amount of work done by `update_acceleration` scales nonlinearly with the number of bodies, let us example the effect of increasing the number of bodies to 250, but decreasing the final time to 1000s,

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
99.13	1.03	1.03	1001	0.00	0.00	update_acceleration(double*, double*, dou
0.96	1.04	0.01	4004	0.00	0.00	euler(double*, double*, double, int)
0.00	1.04	0.00	2496	0.00	0.00	std::__detail::_Mod<unsigned long, 429496
0.00	1.04	0.00	2496	0.00	0.00	unsigned long std::__detail::_mod<unsig
0.00	1.04	0.00	2492	0.00	0.00	std::__detail::_Mod<unsigned long, 624ul,

Increasing the number of bodies further increases the relative time spent inside of `update_acceleration`. If this program is to be used for large numbers of orbiting bodies, doing any optimisations elsewhere would be trivial compared to the potential optimisations awarded by decreasing the run time in `update_acceleration`. `update_acceleration` could be written in the following, relatively naive way:

```

1 void update_acceleration(double p_X[], double p_Y[], double p_M[], double p_aX
   [], double p_aY[]) {
2     // calculate accelerations
3     // set all accelerations to zero
4     for (int i = 0; i < nBodies; i++) {
5         p_aX[i] = 0;
6         p_aY[i] = 0;
7     }
8
9     // for each pair...
10    for (int i = 0; i < nBodies; i++) {
11        for (int j = 0; j < nBodies; j++) {
12            if (i != j) {
13                // get distance between bodies
14                double dist = sqrt(pow(p_X[i] - p_X[j], 2) + pow(p_Y[i] - p_Y[j]
15                ], 2));
16                // calculate component forces
17                double fX = -G * p_M[i] * p_M[j] * (p_X[i] - p_X[j])/pow(dist,
18                3);
19                double fY = -G * p_M[i] * p_M[j] * (p_Y[i] - p_Y[j])/pow(dist,
20                3);
21                // append component accelerations to arrays
22                p_aX[i] += fX / p_M[i];
23                p_aY[i] += fY / p_M[i];
24            }
25        }
26    }
27 }

```

All of the accelerations are set to 0, then, for each of bodies, the force on the i th body is calculated, and accelerations set accordingly. However, we may apply some key optimisations to greatly reduce the work done:

- given the the force acting on object a due to b is equal to the force acting on object b due to a, we can reduce the for loops from looking at every pair (n^2 comparisons) to just have

each body look at the bodies past it in the index, then update both bodies accelerations ($\frac{n^2-n}{2}$ comparisons)

- a lot of the calculations done here are repeated twice - simply precalculating these should decrease runtime by a factor of roughly 2 for these sections

Implementing these improvements, we are left with

```

1 double update_acceleration(double p_X[], double p_Y[],
2                             double p_M[], double p_aX[],
3                             double p_aY[]) {
4     // updates accelerations and returns total potential energy
5     double PE = 0;
6     // for each pair...
7     double dX, dY, dist, coeff, fX, fY, inv_r_cube;
8     for (int i = 0; i < nBodies; i++) {
9         for (int j = i + 1; j < nBodies; j++) {
10             // get distance between bodies
11             dX = p_X[i] - p_X[j]; dY = p_Y[i] - p_Y[j];
12             dist = sqrt(dX * dX + dY * dY); // cartesian dist
13             // calculate component forces
14             coeff = -G * p_M[i] * p_M[j];
15             inv_r_cube = pow(dist, -3);
16             fX = coeff * dX * inv_r_cube;
17             fY = coeff * dY * inv_r_cube;
18             // append component accelerations to arrays
19             p_aX[i] += fX / p_M[i]; p_aX[j] -= fX / p_M[j];
20             p_aY[i] += fY / p_M[i]; p_aY[j] -= fY / p_M[j];
21             // increase potential energy value
22             PE += coeff / dist;
23         }
24     }
25
26     return PE;
27 }

```

Note that, in this version, the acceleration arrays had been set to 0 at a different point in the code.

The integration methods (Euler's method) step size h is linearly proportional to run time, and the choice of which is very dependant on the situation. Thus, trying to "optimise" the code by increase the step size does not really help the situation (although it does reduce calls to `update_acceleration`). Instead, optimisation flags were trialed, and the relation between number of bodies and run time determined experimentally.

1.4 Verification

Due to the nature of the problem, analytical solutions to initial conditions do not exist. Instead, to "verify" that the code is performing correctly, we confirm that the total energy in the system does not vary over time. Since this is a closed system, the total energy should remain constant with time, and a poor integration method would show increases in energy when the situations discussed above occur. This is calculated as follows

```

1 fabs(fabs(final_TE - initial_TE)/initial_TE)

```

To examine the conditions in which the serial code works, the program was run with varying Δt , and the resulting change in energy recorded. This was repeated for a range of n -body counts. The results are plotted in figure 1.1.




Figure 1.1: Comparison of how the energy varies over a 1000s simulation for 2, 16, 32 and 64 bodies.

Below approximately $\Delta t = 1$, there is essentially no change in energy over the run time (that is, within numerical error from calculating the energy components). Past $\Delta t = 0$, the energy variance fluctuates for each set of trials with more than 2 bodies. Consider that the number of potential collisions (or close passes of two bodies) increases with the number of bodies, and hence accelerations spike more often, and the numerical solution is less accurate, in turn leading to jumps in total energy. This is shown below is a pair of ‘good’ and ‘bad’ situations -

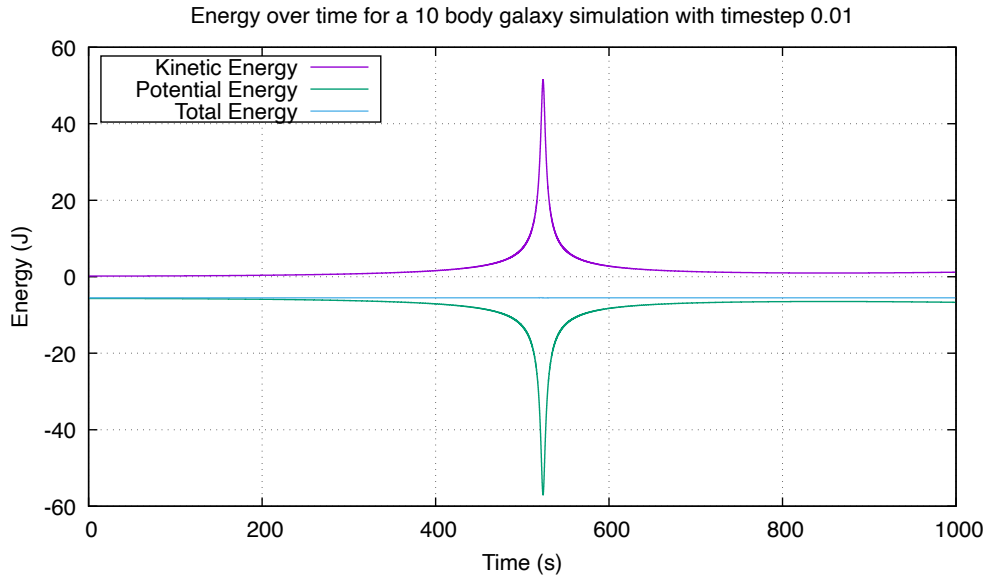


Figure 1.2: Kinetic, potential and total energy over time for a 10 body galaxy simulation, with a small timestep over 1000s

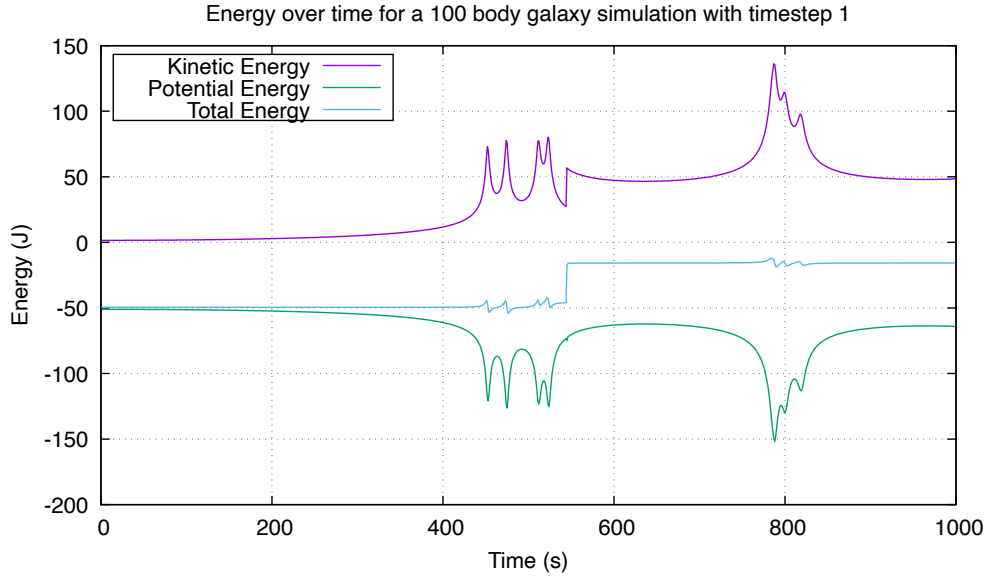


Figure 1.3: Kinetic, potential and total energy over time for a 100 body galaxy simulation, with a large timestep over 1000s

As you can see, the total energy in figure 1.2 is constant, and the trade-off between potential and kinetic energy at a close approach between two planets is smooth, as expected. In figure 1.3, the close approaches between pairs of planets cause discontinuities in the total energy, eventually leading to a significant change in the total energy of the system.

To complement the above results, some basic animations of bodies position over time were produced. Some of these animations are included in the presentation. These showed a couple trends:

- When bodies are far apart, the attractive force is very small, as expected. Conversely, when bodies are closer, the attractive becomes very strong, as expected.
- When two bodies become very close, the attractive force, and hence eventually relative velocity becomes very large. As we are using a fixed step size, this can lead to large jumps in position, perhaps to points that are unphysical. One example that has been observed is that two bodies become close, at one time step that are very close and hence experience a very strong force, and at the next time step are moved very far apart. Since the acceleration at the next time step is relatively low, the two bodies continue on at unphysical speeds.
- In contrast, this is seen very rarely for significantly smaller step sizes.

For simulations henceforth, when we are not examining the relationship between step size or final time and some other parameter, all simulations will be run to a final time of 1000s, with a step size of $\Delta t = 1s$.

1.5 Performance Results

To benchmark the performance of the program, the simulation was run many times with varying parameters, and timed using bash's `time` command. All tests were done with a final time of 1000s, with a time step of 1s. The number of bodies ranged from 10 to 1000. For each number of bodies, the program was run with the optimisation flags -O0, -O1, -O2 and -O3. This was run on `getafix`, and the `goslurm.sh` script is attached below. The flag number, number of bodies, and real time (i.e., wall clock time) were piped to a CSV file, and plotted using gnuplot to produce figure 1.4.




Figure 1.4: Plot of run-times for running a simulation with 1000 time steps, with a range of body counts, and using different compilation flags.

Clearly, the relationship between number of bodies and runtime follows the predicted pattern of $\mathcal{O}(n^2)$. The optimisation flag -O0 performed the worst, and -O2 and -O3 performed similarly well for the best. -O3 took around 75 seconds for 1000 bodies, and -O0 took around 90 seconds for the same parameters, a difference of about 18%.

Given that this model could feasibly be used for a system with many more than n bodies, and the number of time steps here is very small, chosen only as an example, this implementation is currently too slow to be practical.

Chapter 2

Parallel

This chapter attempts to take the serial code created in the previous chapter and decrease the runtime using parallelisation techniques, particularly doing repeated calculations in parallel on a GPU using Cuda. First, there is a discussion of the limitations of paralling this problem, followed by details of the implementation and the verification procedure. Finally, the performance is compared to the serial implementation, particularly how the problem scales with number of bodies.

2.1 Parallelisation Strategies

A key limitation of the n -body problem is that the problem is that each state is dependant on the entirety of the previous state - for each new state, we must first calculate the acceleration of each body, which is dependant on the current location of every state, which can only be calculated at the end of the state. Hence, we cannot, for example, run the `update_acceleration()` function for two states simultaneously, and our parallelisation methods are constrained to only operate on the current state.

First, we shall focus on optimising the `update_acceleration()` function. As shown in the serial gprof, most of the runtime is spent in this function, especially for large numbers of bodies. Thankfully, the calculations for the new state are *only* dependant on the previous state, and as such we can calculate the acceleration for each body without waiting for the other bodies. All of the calculations done in `update_acceleration()` are relatively simple, and so we will have no issue running these using Cuda. It is worth noting, however, the some trickery is required to keep the number of comparisons low.

As discussed earlier, the absolute best way to decrease the runtime of the program is to decrease the number of comparisons made in `update_acceleration()`. This is done by only looking at unique pairs of bodies, rather than all bodies (that is, once we've looked at the pair (1,10), we don't need to look at (10,1)). This was done using careful nested for loops, but for Cuda we'd rather do as little work on each thread at once. Thus, we precalculate the list of pairs of bodies to compare (are relatively instant process), and then use the index of the current Cuda thread to get the current pair to compare.

Beyond this, the implementation is relatively simple, keeping in mind to copy information to and from the device and host as infrequently as possible to save time.

2.2 CUDA Implementation

The implementation of Cuda to parallise the code is fairly straightforward for most of the code - the only functions affected are `main()`, `simulation()`, `time_step()` and, most importantly, `update_acceleration()`. Most of the changes are just a matter of allocating and freeing memory, and copying memory back and forth when needed. Three key changes, however: first, in `main()`,

```

1 // generate (i,j) pairs for body comparisons
2 int pairs = (nBodies * nBodies - nBodies) / 2;
3 int *Is = new int[pairs];
4 int *Js = new int[pairs];
5 int ind = 0;
6 for (int i = 0; i < nBodies; i++) {
7     for (int j = i + 1; j < nBodies; j++) {
8         Is[ind] = i; Js[ind] = j;
9         ind++;
10    }
11 }

```

This generates the list of (i, j) pairs, so that the same optimisation that significantly reduces the working set can still be used when working in parallel.

Next, in `time_step()`,

```

1 // recalculate acceleration arrays
2 update_acceleration<<<pairs / 32 + 1, 32>>>(p_X_device, p_Y_device,
3                                             p_M_device, p_aX_device,
4                                             p_aY_device, Is, Js,
5                                             PE_device, pairs);

```

The call to `update_acceleration()` is careful to minimize the number of blocks and maximize the number of threads. There will be, at most, 31 excess threads initialized.

Finally, in `update_acceleration()`,

```

1 __global__ void update_acceleration(double p_X[], double p_Y[],
2                                     double p_M[], double p_aX[],
3                                     double p_aY[], int Is[], int Js[],
4                                     double *PE_device, int pairs) {
5     // updates accelerations and returns total potential energy
6     // parameters:
7     //   p_X, p_Y: (x,y) values for each body
8     //   p_vX, p_vY: (x,y) velocities for each body
9     //   p_aX, p_aY: (x,y) accelerations for each body
10    //   p_M: mass of each body
11    //   Is, Js: form the (i,j) pairs that need to be calculated for each body
12    //   pairs: integer number of pairs to calculate accelerations for
13    //   PE_device: pointer to potential energy sum on GPU
14    // output:
15    //   no output - overwrites parameter arrays
16    const int ind = blockIdx.x * blockDim.x + threadIdx.x;
17
18    // for each pair...
19    if (ind < pairs) {
20        int i = Is[ind]; int j = Js[ind];
21
22        [sic]
23
24        atomicAdd(PE_device, PE);
25    }
26 }

```

Alongside the earlier additions to code, the lines where the index is calculated to find the (i, j) pair allows the function to be split across as many threads as possible, and sending an even amount of work to each thread. If each thread dealt with a single body, threads would either be doing redundant calculations or the certain threads would have significantly more work than others, leading to many threads waiting.

The overall process by which the code works has not changed, so we should see very similar results to serial in the verification procedure.

2.2.1 Profiling

To profile, we'll run the code once with very few bodies and a long simulation time (20 bodies, timestep 1s, final time 100,000s):

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ns/call	total ns/call	name
37.50	0.03	0.03	100000	300.01	300.01	time_step(double*, double*, double*, double*)
37.50	0.06	0.03				cuda::getCudaError(cudaError_enum)
12.50	0.07	0.01	100000	100.00	100.00	kinetic_energy(double*, double*, double*)
12.50	0.08	0.01				cuda::cuosCondCreateWithSharedFlag(pthread_t*)
0.00	0.08	0.00	100000	0.00	0.00	write_state(double*, double*, double, double)
...						

We see that the time is most distributed between the `time_step()` function and `kinet_energy()` function, and some function used for dealing with memory on the device. This suggests that the CUDA code adds a significant overhead for very few bodies! To compare, we will run the simulation again with 256 bodies, timestep of 1s, and final time of 10,000s,

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
70.00	0.07	0.07	20000	3.50	3.50	time_step(double*, double*, double*, double*)
10.00	0.08	0.01	10000	1.00	1.00	__device_stub__Z19update_accelerationPdS_S
10.00	0.09	0.01				write_header(std::basic_ofstream<char, std::
10.00	0.10	0.01				simulate(double*, double*, double*, double*)

Here, the bulk of the time is still spent in the `time_step` function, but now a significant portion is also spent in `update_acceleration`. Increasing the number of bodies has a very small effect on the time profile of the program - this suggests that the parallelisation techniques have been successful.

2.3 Verification

To verify that the Cuda implementation is performing correctly, the same tests as were performed for serial are used. The system is closed, so we expect to see energy conserved, and hence the y axis to stay near 0,

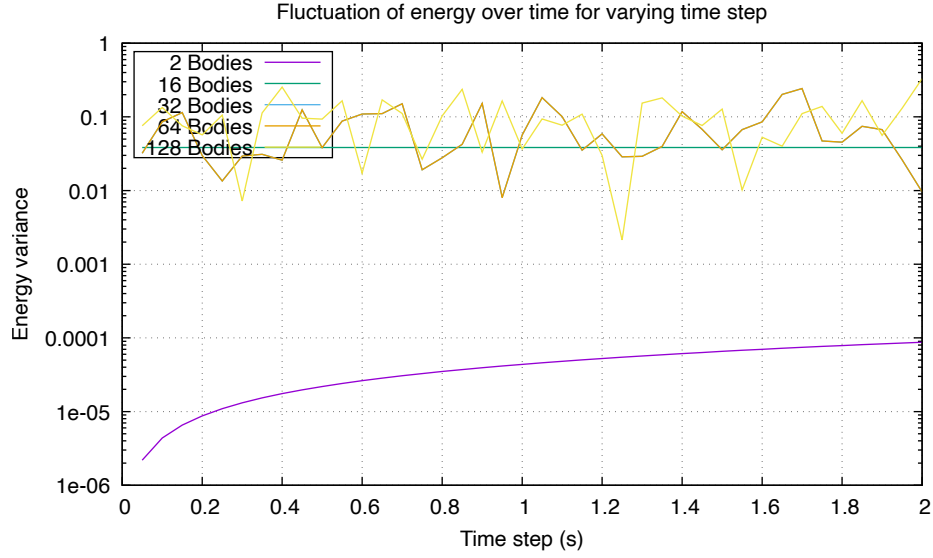


Figure 2.1: Comparison of how the energy varies over a 1000s simulation for 2, 16, 32, 64 and 128 bodies.

Clearly, the step size does not have a significant effect on any of the results - the 2 body system is spaced enough such that the error does not change by any significant amount. The 16 and 32 body systems perform similarly to the serial test, showing that step sizes under $\Delta t = 2$ are suitable. Past this, significant noise is seen in the errors of the 64 and 128 body systems, noticeably more severe in the 128 body system. This is as expected - as the number of bodies increases, so does the density of the bodies and the likelihood of “collisions” (or, at least, particles jumping very close together then far apart), and hence we see poorer performance here.

Next, we examine a “good” and “bad” simulation,

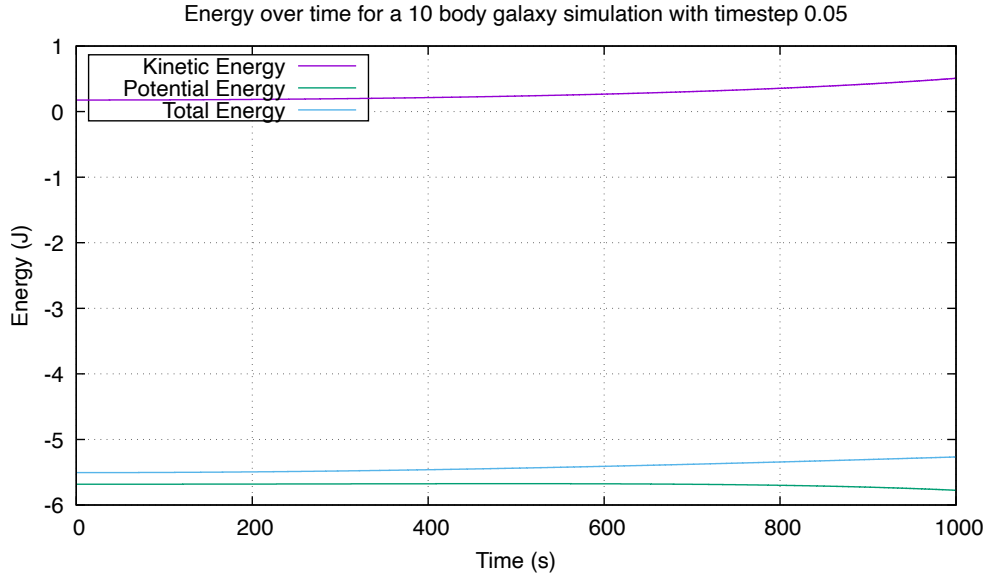


Figure 2.2: Kinetic, potential and total energy over time for a 10 body galaxy simulation, with a small timestep over 1000s




Figure 2.3: Kinetic, potential and total energy over time for a 500 body galaxy simulation, with a large timestep over 1000s

In both of the above scenarios, we slight change in total energy of the run time a constant total energy, a change of approximately 3% for the good situation and approximately 15% for the poor scenario, with a discontinuity near the end of the simulation. This is of course concerning, but considering the change is being observed, in the case of the good scenario, over 20 thousand iterations, I believe this to be an acceptable change. The more drastic change in the poor scenario may be attributed to the significantly larger timestep.

2.4 Experimentation

Finally, we wish to measure the time performance of the paralised code, particularly in comparison to the serial code. To do this, we will run tests with a time step of $\Delta t = 0.01s$, to a final time of 100s, first purely for the Cuda code to determine the best compile flag to use, and then with the serial code to determine which scales better and quantify the difference.




Figure 2.4: Effect of number of bodies of run time of parallel code, with varying optimisation flag

The plot of runtime against problem size in figure 2.4 is quite promising - in the best case, simulating 1000 bodies took only 15s, whereas it took closer to 90s in the serial case! Otherwise, we can observe that the choice of flag has a significant effect on the result. The runtime with optimisation flag -O3 is very consistent. -O2 follows the same trend closely, aside from a jump at 300 and 800 bodies. Flags -O0 and -O1 are relatively ‘noisy’ in comparison, providing wildly varying run time - in some case, much faster than the other flags, however typically worse. As we are most interested in how the code performs as the problem size becomes large, -O3 will be used henceforth.

Now, let us directly compare the performance of the parallel and serial code. We will produce 2 plots - in the first, we plot the run time as a function of the number of bodies. Then, we plot the run time as a function of the final time of the simulation, for a range of different numbers of bodies.

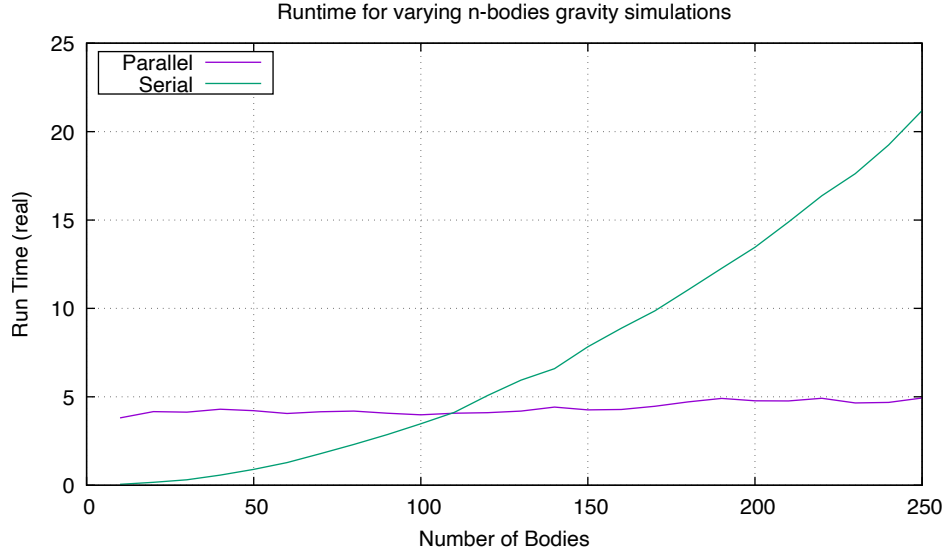


Figure 2.5: Comparison of the run time of the parallel and serial code, for fixed final time and step size (1000s and 1s), and varying number of bodies.

Figure 2.5 shows the run time of the serial code to grow faster than linearly as the number of bodies is increased, and the run time of the parallel code increases only very slightly. Interestingly, for small body counts, below 100, the serial code is faster - this is likely due to the significant overhead associated with running the parallel code (much more initialisation of memory, copying memory around, etc). Clearly, for simulations with many bodies, the overhead is negligible compared to the speed increase, and hence parallel becomes the faster code.

Finally, to compare the effect of changing run time, let's set 3 different body counts: 32, 128, and 256. We should see that the serial code performs better than the parallel code at 32 bodies, similarly at 128, and worse at 256. This is plotted in figure 2.6. As expected, the serial code is faster initially, however becomes slower than the parallel implementation as the final time increases. The rate at which this occurs is proportional to the number of bodies being simulated.

In figure 2.6, the parallel code appears to run in constant time relative to both number of bodies and the final time - in actual fact, it is more likely that the parallel code is so much faster in the `update_acceleration()` function as compared to the serial code that when the two programs are directly compared the parallel code appears to run `update_acceleration()` instantly.

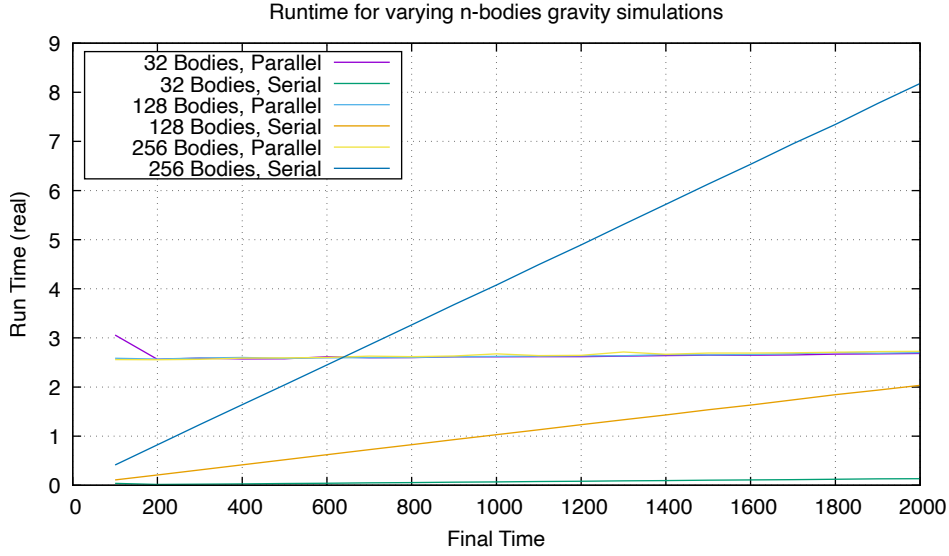


Figure 2.6: Plot of the relationship between run time and final time of simulation, compared between the parallel and serial implementations for a range of body counts.

2.5 Conclusion

This report aimed to propose a naive and later optimised solution to the n -body problem, where some number of bodies experience a gravity force on one another. The first chapter, where the serial implementation was discussed, showed how the amount of work done may be greatly decreased with some careful precalculation and choice of nested for loops. The profile analysis of this code showed that most of the run time was spent calculating the acceleration on each body, so the second chapter discusses how to optimise this function by running it in parallel on a GPU, using CUDA.

Although the opportunities for parallelisation of the simulation are limited - each state requires the entire previous state to be calculated - implementing the simulation on CUDA is successful at providing a very significant speed increase, particularly when simulating scenarios with a large number of bodies and/or a distant final time, despite the overhead that is introduced. This speed increase was so significant that the function to update accelerations was no longer the function where most of the run time was spent. Hence, with the improvements made, the code is now suitable to simulate systems with very many bodies (orders of magnitude than feasible with the naive solution).

Although the optimisation was successful, this report is not without fault or limitations. A more careful and realistic study into the propagation of bodies over time would have strengthened the verification sections - for example, comparing to a known system, like the Milky way, or the analytical solution to a 2/3-body problem. Also, the parallel code was only tested for scenarios that are feasible to test alongside the serial code, due to the serial code being prohibitively slow for large numbers of bodies. Finally, the use of Euler's method for iterating the velocities and position was chosen to minimise the number of calls to the `update_acceleration()` function, with the trade off of accuracy. Since the parallelised code has decreased the run time of `update_acceleration()` so far, a method such as kick-drift-kick or verlet integration (both of which would be required to make more calls to the `update_acceleration()` function) should be implemented and the effect on the verification studied.

Bibliography

- [1] Randall D. Knight. *Physics for Scientists and Engineers*. Pearson, 2017.

Chapter 3

Appendix

3.1 Serial Bash Scripts

3.1.1 Serial Validation

Shell script to submit to sbatch and produce datafiles:

```
1 #!/bin/bash -l
2 #SBATCH --job-name=SerialValidation
3 ###SBATCH --odelist=smp-6-3
4 #SBATCH --nodes=1
5 #SBATCH --ntasks-per-node=1
6 #SBATCH --cpus-per-task=1
7 #SBATCH --mem-per-cpu=2G # memory (MB)
8 #SBATCH --time=0-00:30 # time (D-HH:MM)
9
10 FINAL_TIME=1000
11
12 make serial
13
14 # save example run to file that probably doesn't break
15 build/main_serial --nBodies 10 --finalTime ${FINAL_TIME} --timeStep 0.01 --
    filename output/serial_valid_good.dat
16 # save example run to file that probably does break
17 build/main_serial --nBodies 100 --finalTime ${FINAL_TIME} --timeStep 1 --
    filename output/serial_valid_bad.dat
18
19 echo "dt_2body_16body_32body_64body" > output/serial_valid.csv
20 for dt in $(seq 0.01 0.01 2); do
21     err1=$( { build/main_serial --nBodies 2 --finalTime ${FINAL_TIME} --
        timeStep ${dt} --energyDiff; } 2>&1 )
22     err2=$( { build/main_serial --nBodies 16 --finalTime ${FINAL_TIME} --
        timeStep ${dt} --energyDiff; } 2>&1 )
23     err3=$( { build/main_serial --nBodies 32 --finalTime ${FINAL_TIME} --
        timeStep ${dt} --energyDiff; } 2>&1 )
24     err4=$( { build/main_serial --nBodies 64 --finalTime ${FINAL_TIME} --
        timeStep ${dt} --energyDiff; } 2>&1 )
25     echo "${dt}_${err1}_${err2}_${err3}_${err4}" >> output/serial_valid.csv
26 done
```

Plotting scripts: for the key plot,

```
1 set terminal pdf
```

```

2 set output "serial_valid.pdf"
3
4 set xlabel "Time_step(s)"
5 set ylabel "Energy_variance"
6 set title "Fluctuation_of_energy_over_time_for_varying_time_step"
7 set grid
8
9 set key left top
10 set key box
11
12 plot 'serial_valid.csv' every ::1 using 1:2 with lines title "2_Bodies",\
13      'serial_valid.csv' every ::1 using 1:3 with lines title "16_Bodies",\
14      'serial_valid.csv' every ::1 using 1:4 with lines title "32_Bodies",\
15      'serial_valid.csv' every ::1 using 1:4 with lines title "64_Bodies"

```

For the ‘good’ scenario,

```

1 set terminal pdf
2 set output "serial_valid_good.pdf"
3
4 set xlabel "Time(s)"
5 set ylabel "Energy(J)"
6 set title "Energy_over_time_for_a_10_body_galaxy_simulation_with_timestep_0.01"
7 set grid
8
9 set key left top
10 set key box
11
12 plot 'serial_valid_good.dat' every ::1 using 1:2 with lines title "Kinetic_
    Energy",\
13      'serial_valid_good.dat' every ::1 using 1:3 with lines title "Potential_
    Energy",\
14      'serial_valid_good.dat' every ::1 using 1:4 with lines title "Total_Energy
    "

```

For the ‘bad’ scenario,

```

1 set terminal pdf
2 set output "serial_valid_bad.pdf"
3
4 set xlabel "Time(s)"
5 set ylabel "Energy(J)"
6 set title "Energy_over_time_for_a_100_body_galaxy_simulation_with_timestep_1"
7 set grid
8
9 set key left top
10 set key box
11
12 plot 'serial_valid_bad.dat' every ::1 using 1:2 with lines title "Kinetic_
    Energy",\
13      'serial_valid_bad.dat' every ::1 using 1:3 with lines title "Potential_
    Energy",\
14      'serial_valid_bad.dat' every ::1 using 1:4 with lines title "Total_Energy"

```

3.1.2 Parallel Validations

Shell script to submit to sbatch and produce datafiles:

```

1 #!/bin/bash -l
2 #SBATCH --job-name=ParallelValidation
3 #SBATCH --nodes=1
4 #SBATCH --ntasks=1
5 #SBATCH --ntasks-per-node=1
6 #SBATCH --cpus-per-task=1
7 #SBATCH --partition=gpu
8 #SBATCH --gres=gpu:1
9 #SBATCH --mem-per-cpu=8G # memory (MB)
10 #SBATCH --time=0-00:30 # time (D-HH:MM)
11
12 FINAL_TIME=1000
13
14 make cuda
15
16 module load gnu
17 module load cuda
18
19 # save example run to file that probably doesn't break
20 build/main_cuda --nBodies 10 --finalTime ${FINAL_TIME} --timeStep 0.05 --
    filename output/cuda_valid_good.dat
21 # save example run to file that probably does break
22 build/main_cuda --nBodies 500 --finalTime ${FINAL_TIME} --timeStep 1 --filename
    output/cuda_valid_bad.dat
23
24 echo "dt_2body_16body_32body_64body_128body" > output/cuda_valid.csv
25 for dt in $(seq 0.05 0.05 2); do
26     err1=$( { build/main_cuda --nBodies 2 --finalTime ${FINAL_TIME} --timeStep
        ${dt} --energyDiff; } 2>&1 )
27     err2=$( { build/main_cuda --nBodies 16 --finalTime ${FINAL_TIME} --timeStep
        ${dt} --energyDiff; } 2>&1 )
28     err3=$( { build/main_cuda --nBodies 32 --finalTime ${FINAL_TIME} --timeStep
        ${dt} --energyDiff; } 2>&1 )
29     err4=$( { build/main_cuda --nBodies 64 --finalTime ${FINAL_TIME} --timeStep
        ${dt} --energyDiff; } 2>&1 )
30     err5=$( { build/main_cuda --nBodies 128 --finalTime ${FINAL_TIME} --
        timeStep ${dt} --energyDiff; } 2>&1 )
31     echo "${dt}_${err1}_${err2}_${err3}_${err4}_${err5}" >> output/cuda_valid.
        csv
32 done

```

Plotting scripts same as above with relevant adjustments.

3.2 Performance Analysis

3.2.1 Serial Flags

```
1 #!/bin/bash -l
2 #SBATCH --job-name=GalaxySimTS
3 ###SBATCH --odelist=smp-6-3
4 #SBATCH --nodes =1
5 #SBATCH --ntasks-per-node=1
6 #SBATCH --cpus-per-task=1
7 #SBATCH --mem-per-cpu=1G # memory (MB)
8 #SBATCH --time=0-00:01 # time (D-HH:MM)
9
10 export TIMEFORMAT='%R,%U,%S'
11 N_BODIES=$(seq 2 10 100)
12 # N_BODIES=$(seq 2 10 1000)
13 FLAGS=(0 1 2 3)
14 FINAL_TIME=100
15 TIME_STEP=0.1
16
17 echo "flag,nBodies,rtime,utime,stime" > out.csv
18
19 for flag in ${FLAGS[@]}; do
20     g++ -std=c++11 -O${flag} -Wall main.cpp -o main
21     for n in ${N_BODIES[@]}; do
22         t=$( { time ./main --nBodies ${n} --finalTime ${FINAL_TIME} --
23             timeStep ${TIME_STEP} --noSummary; } 2>&1 )
24         echo "${flag},${n},${t}" >> out.csv
25     done
26 done
```

3.2.2 Parallel Flags

```
1 #!/bin/bash -l
2 #SBATCH --job-name=ParallelValidation
3 #SBATCH --nodes=1
4 #SBATCH --ntasks=1
5 #SBATCH --ntasks-per-node=1
6 #SBATCH --cpus-per-task=1
7 #SBATCH --partition=gpu
8 #SBATCH --gres=gpu:1
9 #SBATCH --mem-per-cpu=8G # memory (MB)
10 #SBATCH --time=0-02:00 # time (D-HH:MM)
11
12 export TIMEFORMAT='%R,%U,%S'
13 # N_BODIES=$(seq 2 20 100)
14 N_BODIES=$(seq 2 10 1000)
15 FLAGS=(2 3)
16 FINAL_TIME=100
17 TIME_STEP=0.01
18
19 module load gnu
20 module load cuda
21
22 echo "flag,nBodies,rtime,utime,stime" > output/cuda_flag_parttwo.csv
23 echo "Starting loop"
```

```

24
25 for flag in ${FLAGS[@]}; do
26     make cuda OLEVEL=-O${flag}
27     for n in ${N_BODIES[@]}; do
28         t=$( { time build/main_cuda --nBodies ${n} --finalTime ${FINAL_TIME}
                --timeStep ${TIME_STEP} --noSummary; } 2>&1 )
29         echo "${flag},${n},${t}" >> output/cuda_flag_parttwo.csv
30     done
31     echo "flag_${flag}_done!"
32 done

```

3.2.3 Misc

n-body comparison

Final time comparison

```

1  #!/bin/bash -l
2  #SBATCH --job-name=ParallelVsSerial
3  #SBATCH --nodes=1
4  #SBATCH --ntasks=1
5  #SBATCH --ntasks-per-node=1
6  #SBATCH --cpus-per-task=1
7  #SBATCH --partition=gpu
8  #SBATCH --gres=gpu:1
9  #SBATCH --mem-per-cpu=8G # memory (MB)
10 #SBATCH --time=0-02:00 # time (D-HH:MM)
11
12 export TIMEFORMAT='%R,%U,%S'
13 N_BODIES=(32, 128, 256)
14 FINAL_TIMES=$(seq 100 200 5000)
15 TIME_STEP=1
16
17 module load gnu
18 module load cuda
19
20 make all
21
22 echo "nBodies,finalTime,rttime_para,utime,stime,rttime_seri,utime,stime" > output
    /run_time_comp.csv
23 echo "Starting_for_loop"
24
25 for n in ${N_BODIES[@]}; do
26     for FINAL_TIME in ${FINAL_TIMES[@]}; do
27         t_para=$( { time build/main_cuda --nBodies ${n} --finalTime ${FINAL_TIME}
                    } --timeStep ${TIME_STEP}; } 2>&1 )
28         t_seri=$( { time build/main_serial --nBodies ${n} --finalTime ${
                    FINAL_TIME} --timeStep ${TIME_STEP}; } 2>&1 )
29         echo "${n},${FINAL_TIME},${t_para},${t_seri}" >> output/run_time_comp.
                    csv
30     done
31     echo "\n\n"
32 done
33 echo "done!"

```


3.3 Main Source Code

3.3.1 Serial

```
1  /*
2  Galaxy Simulation - COSC3500
3  Tom Stephen, 45811449
4
5  Program to simulate gravity interaction between N bodies.
6
7  This is the *serial* implementation.
8
9  See README.md for instructions on usage.
10 */
11
12 #include <cstdlib>
13 #include <iomanip>
14 #include <iostream>
15 #include <math.h>
16 #include <cstring>
17 #include <random>
18 #include <chrono>
19 #include <ctime>
20 #include <fstream>
21
22 #define G 1
23 #define PI 3.1415
24
25 int nBodies = 2;
26 double tf = 1000;
27 double h = 1.0;
28 bool energy_diff = false;
29 bool verbose = false;
30 std::string filename;
31
32 #include "system_init.cpp"
33 #include "debugger_tools.cpp"
34 #include "ode_solve.cpp"
35 #include "helper_functions.cpp"
36
37 double update_acceleration(double p_X[], double p_Y[], double p_M[], double
    p_aX[], double p_aY[]);
38 double time_step(double p_X[], double p_Y[], double p_vX[], double p_vY[],
    double p_aX[], double p_aY[], double p_M[]);
39 void simulate(double p_X[], double p_Y[], double p_vX[], double p_vY[], double
    p_aX[], double p_aY[], double p_M[], std::string filename);
40
41 int main(int argc, char* argv[]) {
42     // parse command line arguments
43     parse_args(argc, argv);
44
45     // create arrays
46     double *planet_X = new double[nBodies];
47     double *planet_Y = new double[nBodies];
48
```

```

49     double *planet_vX = new double[nBodies];
50     double *planet_vY = new double[nBodies];
51
52     double *planet_aX = new double[nBodies];
53     double *planet_aY = new double[nBodies];
54
55     double *planet_M = new double[nBodies];
56
57     // initialise planets
58     init_N_orbitting_planets(planet_X, planet_Y, planet_vX, planet_vY, planet_M
59                             , nBodies);
60
61     // simulate system
62     simulate(planet_X, planet_Y, planet_vX, planet_vY, planet_aX, planet_aY,
63             planet_M, filename);
64
65     // delete arrays
66     delete [] planet_X;
67     delete [] planet_Y;
68
69     delete [] planet_vX;
70     delete [] planet_vY;
71
72     delete [] planet_aX;
73     delete [] planet_aY;
74
75     delete [] planet_M;
76
77     exit(0);
78 }
79
80 void simulate(double p_X[], double p_Y[],
81              double p_vX[], double p_vY[],
82              double p_aX[], double p_aY[],
83              double p_M[], std::string filename) {
84     // simulates n-body system with gravity forces
85     // parameters:
86     //   p_X, p_Y: (x,y) values for each body
87     //   p_vX, p_vY: (x,y) velocities for each body
88     //   p_aX, p_aY: (x,y) accelerations for each body
89     //   p_M: mass of each body
90     //   filename: name & location to write the output file
91     //
92     // No returns, just output file and printing relative change in total
93     // energy.
94
95     // open output file
96     std::ofstream outfile;
97     outfile.open(filename);
98     // headerline
99     write_header(outfile);
100    double PE, KE, TE = 0.0, initial_TE = 0.0, final_TE;
101    // do simulation until time reaches final time 'tf'
102    for (double time = 0.0; time <= tf; time += h) {
103        // step forwards by 'h'

```

```

101     PE = time_step(p_X, p_Y,
102                   p_vX, p_vY,
103                   p_aX, p_aY,
104                   p_M);
105     // calculate energies
106     KE = kinetic_energy(p_vX, p_vY, p_M);
107     TE = PE + KE;
108     if (time == 0.0) {
109         initial_TE = TE;
110     }
111     // output current state
112     write_state(p_X, p_Y, time, KE, PE, TE, outfile);
113 }
114 outfile.close();
115
116 // display relative change in total energy
117 if (energy_diff) {
118     final_TE = TE;
119     std::cout << fabs(fabs(final_TE - initial_TE)/initial_TE) << std::endl;
120 }
121 }
122
123 double time_step(double p_X[], double p_Y[], double p_vX[], double p_vY[],
124                 double p_aX[], double p_aY[], double p_M[]) {
125     // Function to progress simulation forwards by 1 timestep.
126     // This requires calculating the acceleration on each body,
127     // then iterating positions and velocities accordingly.
128     // parameters:
129     //   p_X, p_Y: (x,y) values for each body
130     //   p_vX, p_vY: (x,y) velocities for each body
131     //   p_aX, p_aY: (x,y) accelerations for each body
132     //   p_M: mass of each body
133     // output:
134     //   PE: gravitational potential energy of system
135     //   overwrites p_*[] parameters
136
137     // set all accelerations to zero
138     for (int i = 0; i < nBodies; i++) {
139         p_aX[i] = 0;
140         p_aY[i] = 0;
141     }
142
143     // recalculate acceleration arrays
144     double PE = update_acceleration(p_X, p_Y, p_M, p_aX, p_aY);
145
146     // update velocities
147     euler(p_vX, p_aX, h, nBodies);
148     euler(p_vY, p_aY, h, nBodies);
149
150     // update positions
151     euler(p_X, p_vX, h, nBodies);
152     euler(p_Y, p_vY, h, nBodies);
153
154     return PE;
155 }

```

```

155 double update_acceleration(double p_X[], double p_Y[],
156                             double p_M[], double p_aX[],
157                             double p_aY[]) {
158     // updates accelerations and returns total potential energy
159     // parameters:
160     //   p_X, p_Y: (x,y) values for each body
161     //   p_vX, p_vY: (x,y) velocities for each body
162     //   p_aX, p_aY: (x,y) accelerations for each body
163     //   p_M: mass of each body
164     //   PE: gravitational potential energy of system
165     //   overwrites parameter arrays
166
167     // updates accelerations and returns total potential energy
168     double PE = 0;
169     // for each pair...
170     double dX, dY, dist, coeff, fX, fY, inv_r_cube;
171     for (int i = 0; i < nBodies; i++) {
172         for (int j = i + 1; j < nBodies; j++) {
173             // get distance between bodies
174             dX = p_X[i] - p_X[j]; dY = p_Y[i] - p_Y[j];
175             dist = sqrt(dX * dX + dY * dY); // cartesian dist
176             // calculate component forces
177             coeff = -G * p_M[i] * p_M[j];
178             inv_r_cube = pow(dist, -3);
179             fX = coeff * dX * inv_r_cube;
180             fY = coeff * dY * inv_r_cube;
181             // append component accelerations to arrays
182             p_aX[i] += fX / p_M[i]; p_aX[j] -= fX / p_M[j];
183             p_aY[i] += fY / p_M[i]; p_aY[j] -= fY / p_M[j];
184             // increase potential energy value
185             PE += coeff / dist;
186         }
187     }
188
189     return PE;
190 }

```

3.3.2 Parallel

```

1  /*
2  Galaxy Simulation - COSC3500
3  Tom Stephen, 45811449
4
5  Program to simulate gravity interaction between N bodies.
6
7  This is the *parallel* implementation.
8
9  See README.md for instructions on usage.
10 */
11
12 #include <cstdlib>
13 #include <iomanip>
14 #include <iostream>
15 #include <math.h>
16 #include <cstring>
17 #include <random>

```

```

18 #include <chrono>
19 #include <ctime>
20 #include <fstream>
21
22 #define G 1
23 #define PI 3.1415
24
25 int nBodies = 2;
26 double tf = 1000;
27 double h = 1.0;
28 bool energy_diff = false;
29 bool verbose = false;
30 std::string filename;
31
32 #include "system_init.cpp"
33 #include "debugger_tools.cpp"
34 #include "ode_solve.cpp"
35 #include "helper_functions.cpp"
36
37 void checkError(cudaError_t e);
38 __global__ void update_acceleration(double p_X[], double p_Y[], double p_M[],
    double p_aX[], double p_aY[], int Is[], int Js[], double *PE_device, int
    pairs);
39 void time_step(double p_X[], double p_Y[], double p_vX[], double p_vY[], double
    p_aX[], double p_aY[], double p_M[], double p_X_device[], double
    p_Y_device[], double p_aX_device[], double p_aY_device[], int Is[], int Js
    [], int pairs, double *PE_device);
40 void simulate(double p_X[], double p_Y[], double p_vX[], double p_vY[], double
    p_aX[], double p_aY[], double p_M[], std::string filename, double
    p_X_device[], double p_Y_device[], double p_aX_device[], double p_aY_device
    [], double p_M_device[], int Is[], int Js[], int pairs, double *PE_device);
41
42 int main(int argc, char* argv[]) {
43     // parse command line arguments
44     parse_args(argc, argv);
45
46     // create arrays on host
47     double *planet_X = new double[nBodies];
48     double *planet_Y = new double[nBodies];
49
50     double *planet_vX = new double[nBodies];
51     double *planet_vY = new double[nBodies];
52
53     double *planet_aX = new double[nBodies];
54     double *planet_aY = new double[nBodies];
55
56     double *planet_M = new double[nBodies];
57
58     // ...and on device
59     double *planet_X_device;
60     double *planet_Y_device;
61     double *planet_aX_device;
62     double *planet_aY_device;
63     double *planet_M_device;
64

```

```

65     double *PE_device;
66
67     checkError(cudaMalloc(&planet_X_device, nBodies * sizeof(double)));
68     checkError(cudaMalloc(&planet_Y_device, nBodies * sizeof(double)));
69     checkError(cudaMalloc(&planet_aX_device, nBodies * sizeof(double)));
70     checkError(cudaMalloc(&planet_aY_device, nBodies * sizeof(double)));
71     checkError(cudaMalloc(&planet_M_device, nBodies * sizeof(double)));
72     checkError(cudaMalloc(&PE_device, sizeof(double)));
73
74     // initialise planets
75     init_N_orbitting_planets(planet_X, planet_Y, planet_vX, planet_vY, planet_M
76         , nBodies);
77     // copy initialised states to device
78     checkError(cudaMemcpy(planet_X_device, planet_X, nBodies * sizeof(double),
79         cudaMemcpyHostToDevice));
80     checkError(cudaMemcpy(planet_Y_device, planet_Y, nBodies * sizeof(double),
81         cudaMemcpyHostToDevice));
82     checkError(cudaMemcpy(planet_M_device, planet_M, nBodies * sizeof(double),
83         cudaMemcpyHostToDevice));
84
85     // generate (i,j) pairs for body comparisons
86     int pairs = (nBodies * nBodies - nBodies) / 2;
87     int *Is = new int[pairs];
88     int *Js = new int[pairs];
89     int ind = 0;
90     for (int i = 0; i < nBodies; i++) {
91         for (int j = i + 1; j < nBodies; j++) {
92             Is[ind] = i; Js[ind] = j;
93             ind++;
94         }
95     }
96     // ...and store these pairs on the device
97     int *Is_device;
98     int *Js_device;
99     checkError(cudaMalloc(&Is_device, pairs * sizeof(int)));
100    checkError(cudaMalloc(&Js_device, pairs * sizeof(int)));
101    checkError(cudaMemcpy(Is_device, Is, pairs * sizeof(int),
102        cudaMemcpyHostToDevice));
103    checkError(cudaMemcpy(Js_device, Js, pairs * sizeof(int),
104        cudaMemcpyHostToDevice));
105
106    // simulate system
107    simulate(planet_X, planet_Y, planet_vX, planet_vY, planet_aX, planet_aY,
108        planet_M, filename, planet_X_device, planet_Y_device, planet_aX_device,
109        planet_aY_device, planet_M_device, Is_device, Js_device, pairs,
110        PE_device);
111
112    // free memory
113    delete [] planet_vX;
114    delete [] planet_vY;
115
116    checkError(cudaFree(planet_X_device));
117    checkError(cudaFree(planet_Y_device));
118    checkError(cudaFree(planet_aX_device));
119    checkError(cudaFree(planet_aY_device));

```

```

111     checkError(cudaFree(planet_M_device));
112     checkError(cudaFree(PE_device));
113
114     exit(0);
115 }
116
117
118 void simulate(double p_X[], double p_Y[],
119             double p_vX[], double p_vY[],
120             double p_aX[], double p_aY[],
121             double p_M[],
122             std::string filename,
123             double p_X_device[], double p_Y_device[],
124             double p_aX_device[], double p_aY_device[],
125             double p_M_device[],
126             int Is[], int Js[], int pairs,
127             double *PE_device) {
128     // simulates n-body system with gravity forces
129     // parameters:
130     //   p_X, p_Y: (x,y) values for each body
131     //   p_vX, p_vY: (x,y) velocities for each body
132     //   p_aX, p_aY: (x,y) accelerations for each body
133     //   p_M: mass of each body
134     //   filename: name & location to write the output file
135     //   p_*_device: as above, but stored on the GPU
136     //   Is, Js: form the (i,j) pairs that need to be calculated for each body
137     //   PE_device: pointer to potential energy sum on GPU
138     //
139     // No returns, just output file and printing relative change in total
140     // energy.
141
142     // open output file
143     std::ofstream outfile;
144     outfile.open(filename);
145     // headerline
146     write_header(outfile);
147
148     double PE, KE, TE, initial_TE, final_TE;
149     // do simulation until time reaches final time 'tf'
150     for (double time = 0.0; time < tf; time += h) {
151         // step forwards by 'h'
152         time_step(p_X, p_Y, p_vX, p_vY, p_aX, p_aY,
153                 p_M_device, p_X_device, p_Y_device, p_aX_device, p_aY_device,
154                 Is, Js, pairs, PE_device);
155         // calculate energies
156         KE = kinetic_energy(p_vX, p_vY, p_M);
157         checkError(cudaMemcpy(&PE, PE_device, sizeof(double),
158                             cudaMemcpyDeviceToHost));
159         TE = PE + KE;
160         if (time == 0.0) {
161             initial_TE = TE;
162         }
163         // output current state
164         write_state(p_X, p_Y, time, KE, PE, TE, outfile);
165     }

```

```

164     outfile.close();
165
166     // display relative change in total energy
167     if (energy_diff) {
168         final_TE = TE;
169         std::cout << fabs(fabs(final_TE - initial_TE)/initial_TE) << std::endl;
170     }
171 }
172
173 void time_step(double p_X[], double p_Y[], double p_vX[], double p_vY[], double
    p_aX[], double p_aY[], double p_M_device[], double p_X_device[], double
    p_Y_device[], double p_aX_device[], double p_aY_device[], int Is[], int Js
    [], int pairs, double *PE_device) {
174     // Function to progress simulation forwards by 1 timestep.
175     // This requires calculating the acceleration on each body,
176     // then iterating positions and velocities accordingly.
177     // parameters:
178     //     p_X, p_Y: (x,y) values for each body
179     //     p_vX, p_vY: (x,y) velocities for each body
180     //     p_aX, p_aY: (x,y) accelerations for each body
181     //     p_M: mass of each body
182     //     p_*_device: as above, but stored on the GPU
183     //     Is, Js: form the (i,j) pairs that need to be calculated for each body
184     //     PE_device: pointer to potential energy sum on GPU
185     // output:
186     //     no output - overwrites p_*[] parameters
187
188     // set all accelerations to zero
189     for (int i = 0; i < nBodies; i++) {
190         p_aX[i] = 0;
191         p_aY[i] = 0;
192     }
193     checkError(cudaMemcpy(p_aX_device, p_aX, nBodies * sizeof(double),
        cudaMemcpyHostToDevice));
194     checkError(cudaMemcpy(p_aY_device, p_aY, nBodies * sizeof(double),
        cudaMemcpyHostToDevice));
195
196     // set potential energy sum to 0
197     double reset_PE = 0;
198     checkError(cudaMemcpy(PE_device, &reset_PE, sizeof(double),
        cudaMemcpyHostToDevice));
199
200     // recalculate acceleration arrays
201     update_acceleration<<<pairs / 32 + 1, 32>>>(p_X_device, p_Y_device,
        p_M_device, p_aX_device, p_aY_device, Is, Js, PE_device, pairs);
202
203     // copy accelerations back to host
204     checkError(cudaMemcpy(p_aX, p_aX_device, nBodies * sizeof(double),
        cudaMemcpyDeviceToHost));
205     checkError(cudaMemcpy(p_aY, p_aY_device, nBodies * sizeof(double),
        cudaMemcpyDeviceToHost));
206
207     // update velocities
208     euler(p_vX, p_aX, h, nBodies);
209     euler(p_vY, p_aY, h, nBodies);

```



```

210 // update positions
211 euler(p_X, p_vX, h, nBodies);
212 euler(p_Y, p_vY, h, nBodies);
213
214 // copy positions back to device
215 checkError(cudaMemcpy(p_X_device, p_X, nBodies * sizeof(double),
216                      cudaMemcpyHostToDevice));
216 checkError(cudaMemcpy(p_Y_device, p_Y, nBodies * sizeof(double),
217                      cudaMemcpyHostToDevice));
217 }
218
219 __global__ void update_acceleration(double p_X[], double p_Y[],
220                                   double p_M[], double p_aX[],
221                                   double p_aY[], int Is[], int Js[],
222                                   double *PE_device, int pairs) {
223 // updates accelerations and returns total potential energy
224 // parameters:
225 //   p_X, p_Y: (x,y) values for each body
226 //   p_vX, p_vY: (x,y) velocities for each body
227 //   p_aX, p_aY: (x,y) accelerations for each body
228 //   p_M: mass of each body
229 //   Is, Js: form the (i,j) pairs that need to be calculated for each body
230 //   pairs: integer number of pairs to calculate accelerations for
231 //   PE_device: pointer to potential energy sum on GPU
232 // output:
233 //   no output - overwrites parameter arrays
234 const int ind = blockIdx.x * blockDim.x + threadIdx.x;
235
236 // for each pair...
237 if (ind < pairs) {
238     int i = Is[ind]; int j = Js[ind];
239     double dX, dY, dist, coeff, fX, fY, inv_r_cube;
240
241     // get distance between bodies
242     dX = p_X[i] - p_X[j]; dY = p_Y[i] - p_Y[j];
243     dist = sqrt(dX * dX + dY * dY); // cartesian dist
244     // calculate component forces
245     coeff = -G * p_M[i] * p_M[j];
246     inv_r_cube = pow(dist, -3);
247     fX = coeff * dX * inv_r_cube;
248     fY = coeff * dY * inv_r_cube;
249     // append component accelerations to arrays
250     p_aX[i] += fX / p_M[i]; p_aX[j] -= fX / p_M[j];
251     p_aY[i] += fY / p_M[i]; p_aY[j] -= fY / p_M[j];
252     // increase potential energy value
253     double PE = coeff / dist;
254
255     atomicAdd(PE_device, PE);
256 }
257 }
258
259 void checkError(cudaError_t e)
260 {
261     if (e != cudaSuccess)
262     {

```

```

263     std::cerr << "CUDA_error:" << int(e) << ":" << cudaGetErrorString(e)
        << '\n';
264     abort();
265 }
266 }

```

3.3.3 Helper functions

General

```

1 void parse_args(int argc, char* argv[]) {
2     for (int i = 1; i < argc; ++i) {
3         if(!strcmp(argv[i], "-h") || !strcmp(argv[i], "--help")) {
4             std::cout << "Galaxy_Simulation-for-COSC3500,byTomStephen" << std
                ::endl
5                 << std::endl
6                 << "Usage:./main_[arguments]run_some_simulation
                " << std::endl
7                 << "\tor:_vim_[arguments]>_[filename]output_results_to_a
                _data_file" << std::endl
8                 << std::endl
9                 << "Arguments:" << std::endl
10                << "\t--nBodiesSet_number_of_bodies_to_simulate" <<
                std::endl
11                << "\t--finalTimeSet_time_to_simulate_to" << std::endl
12                << "\t--timeStepSet_time_step('h')_for_integration_
                method" << std::endl
13                << "\t--energyDiffShow_change_in_total_energy_across_
                simulation" << std::endl
14                << "\t--timeProvide_the_current_time_to_the_
                program" << std::endl;
15            exit(0);
16        } else if (!strcmp(argv[i], "--nBodies")) {
17            nBodies = atoi(argv[++i]);
18        } else if (!strcmp(argv[i], "--finalTime")) {
19            tf = atof(argv[++i]);
20        } else if (!strcmp(argv[i], "--timeStep")) {
21            h = atof(argv[++i]);
22        } else if (!strcmp(argv[i], "--energyDiff")) {
23            energy_diff = true;
24        } else if (!strcmp(argv[i], "--filename")) {
25            filename = argv[++i];
26        }
27    }
28 }
29
30 void write_state(double p_X[], double p_Y[], double t, double KE, double PE,
    double TE, std::ofstream& output) {
31     // function to write current state to file.
32     // writes current time, component and total energies,
33     // and the (x,y) position of each body.
34     output << t << " " << KE << " " << PE << " " << TE;
35     for (int i = 0; i < nBodies; i++) {
36         output << " " << p_X[i] << " " << p_Y[i];
37     }
38     output << "\n";

```

```

39 }
40
41 void write_header(std::ofstream& output) {
42     // function to write the header line out the output file
43     output << "t_KE_PE_TE";
44     for (int i = 0; i < nBodies; i++) {
45         output << "p" << i << "x_p" << i << "y";
46     }
47     output << "\n";
48 }
49
50 double kinetic_energy(double p_vX[], double p_vY[], double p_M[]) {
51     // function to calculate the total kinetic energy of the system.
52     // parameters:
53     //     p_vX: x velocities of each body
54     //     p_vY: y velocities of each body
55     //     p_M: mass of each body
56     double KE = 0; // initialise sum
57     double v_sq; // declare v^2
58     for (int i = 0; i < nBodies; i++) {
59         // calculate v^2
60         v_sq = pow(p_vX[i],2) + pow(p_vY[i],2);
61         // increase sum energy
62         KE += 0.5 * p_M[i] * v_sq;
63     }
64     return KE;
65 }

```

Debugger

```

1 #include <cstdlib>
2 #include <iostream>
3
4 void disp_planets(double planet_X[], double planet_Y[], double planet_vX[],
5     double planet_vY[], double planet_M[], int nBodies) {
6     // usage: disp_planets(planet_X, planet_Y, planet_vX, planet_vY, planet_M,
7     // nBodies);
8     for (int i = 0; i < nBodies; i++) {
9         std::cout << "Planet_" << i << std::endl;
10        std::cout << "(x,y)_" << planet_X[i] << "," << planet_Y[i] << ")" <<
11        std::endl;
12        std::cout << "(vx,vy)_" << planet_vX[i] << "," << planet_vY[i] << ")"
13        << std::endl;
14        std::cout << "m_" << planet_M[i] << std::endl;
15        std::cout << std::endl;
16    }
17 }

```

Initialisation

```

1 #include <cstdlib>
2 #include <math.h>
3 #include <random>
4
5 #define PI 3.1415

```

```

6 #define G 1
7
8 void init_N_orbitting_planets(double planet_X[], double planet_Y[], double
  planet_vX[], double planet_vY[], double planet_M[], int nBodies) {
9     // massive star at origin (index 0). n-1 planets orbit it.
10    // Orbital velocity calculated for each, moving in clockwise direction
11    // system parameters
12    double mass = 1;
13    double dist_max = 5000;
14    double dist_min = 500;
15    double star_mass = 1000;
16
17    // set star values
18    planet_X[0] = 0; planet_Y[0] = 0; planet_vX[0] = 0; planet_vY[0] = 0;
19    planet_M[0] = star_mass;
20
21    // setup random number generator
22    int seed = 69;
23    std::mt19937 generator;
24    std::mt19937 new_gen;
25    generator.seed(seed);
26    new_gen.seed(seed + 1);
27    std::uniform_real_distribution<double> distribution(0.0,1.0);
28    std::uniform_real_distribution<double> new_dist(0.25,0.25);
29
30    // for each remaining planet...
31    double dist, angle, orb_speed;
32    for (int i = 1; i < nBodies; i++) {
33        // generate position + velocity for new body
34        planet_M[i] = mass;
35
36        // get random height in [0, dist_bound]
37        dist = dist_min + (dist_max - dist_min) * distribution(generator);
38
39        // get random angle in [0, 2pi]
40        angle = 2 * PI * distribution(generator);
41
42        // get and set x and y to suit
43        planet_X[i] = dist * cos(angle);
44        planet_Y[i] = dist * sin(angle);
45
46        // calculate orbital velocity
47        // orb_speed = sqrt(G * star_mass / dist);
48        // orb_speed = sqrt(G * star_mass / dist) * 0.5;
49        orb_speed = sqrt(G * star_mass / dist) * new_dist(new_gen);
50        // orb_speed = 0;
51
52        // get and set vx and vy to suit
53        planet_vX[i] = orb_speed * cos(angle + PI / 2);
54        planet_vY[i] = orb_speed * sin(angle + PI / 2);
55    }
56 }

```

Iteration

```

1 void euler(double x[], double dX[], double dt, int n) {

```

```
2   for (int i = 0; i < n; i++) {  
3       x[i] += dt * dX[i];  
4   }  
5 }
```